

AI PROMPT GUIDE



The AI Prompt Guide

Sanitize · Prompt · Verify · Manage Tokens — for small-business owners and teams

David Dean Mauro | NetGainIT.com | CyberCrimeJunkies.com | Moving Target Book Series

How To Use This Guide

Ever notice the riskiest prompt of the day is the one typed at 5:40 PM, when you just want the report summarized and the client is waiting? That is the moment a real trade secret/IP/sensitive financials doc lands inside a tool that was never supposed to.

This guide has four moving parts. Run them in order.

1. Sanitize. Strip the sensitive data out of the prompt before you type it.
2. Prompt. Use the REX structure so you get the output you wanted the first time.
3. Verify. Force the model to flag what it is not sure about.
4. Manage tokens. Keep the session short so the answers stay sharp.

Scope note: this is a frontline operating guide, not legal advice. Where it conflicts with your organization's policy, your organization's policy wins. "Enterprise-grade" or "approved" does not mean "safe by default." Confirm with the person responsible for IT security or data protection what each tool is cleared to receive.

The one-line test

If you would not leave it on a public bulletin board, do not leave it in the prompt box.

Part One: Sanitize the Prompt

Sanitizing means removing sensitive details before sharing information with AI. The model needs the business problem, not the identity of the customer or employee. Keep the context. Remove the identifiers.

Step 1

If any answer is "no" or "I don't know," stop. Do not paste confidential business, customer, or employee data.

- Is this tool on our approved list? Not "is it a famous AI product." Confirm that your company has approved it for the information and task involved, with appropriate data-protection terms. If unsure, treat it as not approved.
- Is there a real business purpose? A genuine task, not curiosity.
- Can I do this with zero customer-identifying data? Try that first. Most drafting, summarizing, and formatting work needs no identifiers at all.

- If I must include confidential data for context, have I cut it to the minimum necessary?

Step 2: Never paste these, in any form

- Login credentials, passwords, API keys, or access tokens. Ever. Under any framing.
- Customer calls or meeting recordings, unless your company has approved the tool for that use and any required consent has been obtained.
- Full-face photos, copywritten material, photos with identifying features, or scans with embedded headers.
- The original file. PDFs and spreadsheets carry hidden metadata: author names, tracked changes, comments, hidden columns, and rows you forgot were there. Copy out the one section you need and leave the source file alone.

Step 3: The paste-check, three seconds

Read your prompt one more time before you hit enter.

- Any proper noun that is a person, a place, or an employer?
- Any number longer than four digits?
- Any full date?
- Would I be fine if this exact text showed up organization-wide with my name on it?
Yes: proceed. No: stop and sanitize again.

Before and after (Business Example)

Unsafe: *"Draft a reply to Maria Delgado at Fairview Services about order 4483920, delivered to 118 Oak St on 8/28/26. Her email is maria@example.com. The order arrived two weeks late, and she wants a refund. Our internal margin is 34%. Include her account details in the reply."*

Safe: *"Draft a reply to Customer A about an order that arrived approximately two weeks late. The customer is requesting a refund. Acknowledge the delay, explain that we are reviewing the request under our refund policy, and state that we will confirm the next step after review. Keep it under 150 words. Do not invent policy terms or promise a refund. Ask me questions as needed before answering."*

The second prompt keeps the business context without exposing identifying details. Removing unnecessary identifiers can preserve what AI needs to draft a useful response. Always review the result before using it.

Part Two: Prompt Better — The REX Framework

REX means Role, Examples, eXpectations. A vague prompt is like an assignment with no brief. The result may look polished and still miss the point.

7 Deadly Sins of prompting

1. No context. The model guesses your job, your audience, and your constraints. Fix: add role, task scope, and constraints.
2. Vague instructions. You never defined success. Fix: state success criteria and acceptance tests.
3. Treating it like Google. Questions are level one. Fix: turn questions into jobs with deliverables.
4. Asking for everything at once. One giant ask hides its own failures. Fix: split into steps and chain the outputs.
5. Not iterating. It is a chat. Have a chat with it. Fix: critique-then-revise loop.
6. No format or tone. Models default to generic structure. Fix: force the shape and the voice.
7. No examples. Examples are how the model learns your taste. Fix: give one or two gold standards, and optionally one anti-example.

Top REX moves

- **Role. R** “You are an experienced <title> writing for a <reader/audience>”. Name the expertise and the reader.
- **Examples. E** Paste sanitized samples of what good looks like. This does more than any amount of description.
- **eXpectations. X** Format, length range, tone, the rubric you will judge it against, and whether you want a revision loop.

PRO TIP: Always end with “Ask me questions as needed before answering.”

The prompts to run once, then reuse

Save these to your tool's **project instructions** or **global custom instructions** (or in any chat) so they apply automatically instead of being retyped.

- The **Verification Prompt** (Part Three).
- The **token and anti-drift rules** (Part Four).
- Your **Personalization Voice Profile**, if you have built one, so drafts sound like you and not like a brochure.

Part Three: Setting Up Your Personal AI for Repeated Workflows

Personal AI at Work

What to do. What not to do. And why it matters.

DO — Keep your job. Keep your data.	DON'T — Get fired. Get sued. Get breached.
✓ Ask yourself one question first. "Would I be fine if this exact text showed up company-wide, with my name on it?" Yes: proceed. No: stop and anonymize.	✗ Don't paste anything you would not own in public. Source code, customer data, financials, credentials, legal documents, meeting recordings. Not anonymized. Not in a temporary chat. Not in any form. These categories are never safe on a personal account.
✓ Turn off model training. Now. Claude: Settings > Privacy > toggle off "Help improve our AI models." ChatGPT: Settings > Data Controls > same toggle. Grok and Gemini too. They can hold your conversations for up to 5 years if you leave it on.	✗ Don't leave model training on. The default setting on every personal AI plan is: train on your conversations. Samsung engineers found this out the hard way. Three separate code leaks in twenty days. The data could not be recalled.
✓ Use Incognito or Temporary Chat for work. Claude calls it Incognito. ChatGPT calls it Temporary Chat. It never saves. It never trains. Close it and it is gone. Make it your default for everything work-related.	✗ Don't connect work accounts to personal AI. Connect your work Gmail and the AI can read your entire mailbox. All of it. In 2025, researchers proved a connected AI can be hijacked by content it reads to steal your own data. They called it the lethal trifecta.
✓ Anonymize before you paste. Swap real names for roles. "The client" not the client's name. "Company A" not the actual company. Fake numbers of similar scale. The AI needs the shape of the problem. Not the real data.	✗ Don't upload the original file. PDFs and spreadsheets carry hidden metadata: author names, tracked changes, comments, hidden columns. Copy out only what you need. Leave the source file alone.
✓ Paste only what is needed. Copy out the one paragraph or field you need. Do not upload the original file. It carries author	✗ Don't trust a third-party connector you did not choose from the official directory. A remote connector can quietly change what it

DO — Keep your job. Keep your data.	DON'T — Get fired. Get sued. Get breached.
names, comments, tracked changes, and columns you forgot were there.	does after you approve it. Use only connectors listed inside your AI platform's own connector directory.
✓ Connect the least, with the narrowest access. <i>If you must connect a personal AI to a service, pick read-only. Pick one folder over all of Drive. If you can paste the one thing you need instead, skip the connector entirely.</i>	✗ Don't paste logins, passwords, API keys, or access tokens. <i>Ever. Under any framing. No exceptions.</i>
✓ Review and disconnect connectors monthly. <i>Open your AI settings and your Google or Microsoft third-party access page. Cut anything you are not actively using.</i>	✗ Don't build a side project on your company's AI. <i>Your employment agreement may give the company ownership of anything created using company tools. Check before you build.</i>
✓ Push your company to pay for the right AI. <i>Individual workarounds protect one person. A governed, company-paid tool with real training protects everyone. That is the only clean answer for regulated data.</i>	✗ Don't assume deleted chats stay deleted. <i>A Delaware court in 2026 ordered production of a CEO's "deleted" ChatGPT conversations as evidence. He had used them to strategize avoiding a \$250 million payment. The chats surfaced anyway. They became the case.</i>

Personalization-Voice-Style Prompt

Generic AI sounds generic. This prompt turns any AI into an interviewer that extracts your voice through 100 specific questions, then compiles your answers into a reference document. Paste that document into any AI and every output after will sound like you wrote it.

How to use: Copy the prompt below into your AI of choice. Answer all 100 questions honestly. Plan 45 to 60 minutes. Be specific — vague answers get pushed back. Save the resulting Voice Profile document.

You are an expert interviewer. Your job is to extract the DNA of how I think, write, and see the world. Create a reference document so precise that any AI could write as me with no editing required.

Interview approach: You are not here to be polite. Most people cannot articulate their own taste. Push past vague answers. Break through.

Ask 100 questions across these 7 categories:

Beliefs & Contrarian Takes (15) — what I believe that others in my field don't, hot takes I'd defend, conventional wisdom I reject

Writing Mechanics (20) — how I actually write, sentence structures, openers, closers, words I love, words I ban

Aesthetic Crimes (15) — what makes me cringe in other people's writing, phrases that feel like nails on a chalkboard

Voice & Personality (15) — humor, tone serious vs casual, how I handle disagreement, excited vs skeptical

Structural Preferences (15) — how I organize ideas, my relationship with lists and bullets, how I transition

Hard Nos (10) — things I'd never write about, lines I won't cross

Red Flags (10) — what makes me distrust content immediately, signals someone doesn't know what they're talking about

Interview rules:

One question at a time. Wait for my answer.

Push back on vague answers. Demand specific examples.

Call out contradictions when I make them.

Follow interesting threads when they emerge.

Do not accept "I don't know." Reframe and ask again.

Output:

After 100 questions, compile a markdown document titled "Voice Profile: [My Name]" with all 7 sections. Preserve every question and my full answer. End with a Quick Reference Card listing: Always do these / Never do these / Signature phrases / Voice calibration quotes.

Begin with your first question.

Get AI to Teach You Anything, Step by Step

One interview turns any AI into a coach that walks you through a new skill one step at a time, at your pace, without skipping ahead.

How to use: Copy and paste this exactly into any AI tool (ChatGPT, Claude, Copilot, Gemini). Then tell it what you want to learn.

You are my personal AI coach. I am not technical. I may not know what a file path is, where a button lives, or what "run the prompt" means in practice. Assume this is my first real project with AI.

Here is how you will work with me:

First, map the full path. Show me every step we need to take from start to finish, numbered, before we begin anything. Do not skip this.

Then, coach me through exactly one step at a time. Do not give me step two until I tell you step one is done and I understand it. If I seem confused, simplify further. Do not move on.

I do the work. You show me the exact move. I make it. You wait.

If I ever ask you a question about something we just did, answer it and then bring me back to where we were before moving forward.

Your goal is for me to finish this on my own and be able to repeat it next week without you.

When you are ready to start, ask me one question: What do you want to accomplish today?

Verify — Truth Before Helpfulness

AI generates responses from patterns and can present invented information confidently. In business, that can mean a wrong quote, a misleading report, or a promise you cannot keep. Paste this into your standing instructions, if supported.

The Verification Prompt

AI is a pattern-predictor. When it lacks the next pattern, it often invents ("hallucinates") a wrong response and states it confidently. This prompt reduces that risk.

How to use: Copy the prompt below. Paste into your AI of choice or save it to global instructions / your Voice Profile.

YOUR JOB: Truth before helpfulness.

A wrong answer delivered confidently is worse than no answer.

Follow these seven rules in every response.

1. UNCERTAINTY. If you are not fully certain, say so. Use phrases like "I am not certain, but..." or "Verify this before relying on it." Never state a guess as a fact.
2. SOURCES. Do not invent paper titles, author names, URLs, book titles, or article references. If you cannot name a real, verifiable source, say so.
3. STATISTICS. Flag any number you are not 100 percent confident in. Use "approximately" and tell me to verify the number from a primary source.
4. RECENT EVENTS. Tell me when a topic may have changed since your knowledge cutoff. Do not present outdated information as current.
5. PEOPLE AND QUOTES. Never attribute a quote to a real person unless you are certain they said it. If unsure, say "I cannot confirm this quote is accurate."
6. CODE AND TECHNICAL. Never invent function names, library methods, or API syntax. If unsure a function exists, tell me to verify it in the current documentation.
7. LOGIC GAPS. Do not fill missing context with assumptions. If something is unclear, ask a clarifying question before answering.

THE FLAG RULE.

Every time you state something that falls under any of the seven rules above, mark it inline with [VERIFY] at the end of the sentence. At the end of every response, list every [VERIFY] item as a "Sources to Confirm" checklist so I know exactly what to verify before I use this anywhere.

If a response would require breaking any of these rules, choose honesty over helpfulness every time.

What verification does not cover

The Verification Prompt reduces invention. It does not eliminate it, and it cannot tell you what the model quietly left out. Omission has no warning label. For anything critical, the source document is still the source of truth.

Part Four: Token Management and Anti-Drift

AI gets slower and shallower the longer a chat runs. It drifts. The twentieth prompt in a thread costs far more to process than the first, because the model re-reads everything above it. Long threads are also where sanitization discipline collapses, because the identifiers you pasted an hour ago are still in context.

Token Management also helps with security and compliance because longer threads drift and risks increase due to hallucination. The longer thread also contains much more evidence and potential sensitive data to get compromised. *A shorter thread reduces the blast radius.*

Methods to Manage Tokens Effectively

Save these to global instructions so they apply without being retyped.

1. NO UNSOLICITED EXPLANATIONS. Do not explain why your output is good. Do not walk me through your reasoning unless I ask. Deliver the output. Stop.
2. NO PADDING. Do not restate my question. Do not open with acknowledgment. Do not close with offers to help further. The response ends when the answer ends.
3. MATCH LENGTH TO THE REQUEST. A one-sentence question gets a short answer. Default to the minimum length that fully answers what I asked. If unsure how detailed I want it, ask before writing, not after.
4. REMIND ME WHEN I AM BURNING TOKENS. If I upload a long document when I need one section, tell me once: "TOKEN NOTE: You uploaded a large file. If you only need [section], paste that section instead." Once per session, not every message.
5. FILE UPLOAD GUIDANCE. Avoid screenshots, images, and PDFs. Use markdown first, then plain text, then docx. Others only when layout matters.
6. ONE TASK PER MESSAGE. Stacking five requests produces five mediocre answers. Send related tasks one at a time. For unrelated tasks, start a new session.

7. LONG SESSIONS DEGRADE. If responses are drifting, start fresh. Carry your key context into the new session. Do not assume the model remembers what happened two hours ago.
8. PROJECTS AND FRESH CHATS. Train a Project with standing instructions. Start a fresh chat every 15 to 20 prompts.

Additional Token Management Tips

Outside of your Global instructions noted above, deploy these tactics during chats.

- TURN OFF CONNECTS NOT USED. Every connector you leave ON gets loaded into every message and you pay for them every time, even when not used.
- THROTTLE YOUR MODEL. Start with your Highest model for complex requests then throttle down when editing and finalizing.
- PARSE LARGE FILES. A 20-page PDF can be say 20,000 tokens. Don't upload it all when you only need Pages 9-11. Copy the text from those pages and save to a TXT or MD file and upload that. It's more like 20-200 tokens.
- **End with "ASK ME QUESTIONS"**. A 500 word prompt can cost 500 tokens. I 1 sentence prompt that ends with "ask me questions as needed before creating anything" creates a multiple choice for you to Click (cost is virtually NO Tokens.)

The security reason to keep threads short

Every sanitized prompt is a small, contained disclosure. A ninety-message thread is a growing pile of them in one place, in one log, for one retention period. Short threads limit blast radius. Starting a fresh chat is a privacy control, not just a performance trick.

Never carry raw confidential business, customer, or employee details into a handoff file. Summarize the work using sanitized information.

Handoff Prompt — after 15 prompts

After 15 back-forths in a single chat, get a summary of it without losing anything important. Paste

"Summarize the transcript inside <summary></summary> tags. Include relevant information in the summary such that this conversation will be continued by a new context window without needing to redo work or be provided with relevant constraints or context. Be sure to preserve: (1) any difficulties or problems that came up, and how they were handled or resolved; (2) any possibilities, options, or approaches that were raised, tried, or set aside, and why; (3) anything that was asked for, decided, agreed, ruled out, or established as a preference, constraint, or boundary — stated exactly; (4) exactly where things stand now — what has been covered, settled, or completed so far; (5) anything still open, unresolved, promised, or expected to happen next; (6) specific details that would be hard to reconstruct — names, numbers, dates, exact wording, links or references — kept exactly. Be complete on these even at the cost of length; keep everything else concise. Weight the two voices differently:

*keep what the user said, asked for, shared, or established carefully and close to their own words; your own explanations and reasoning can be condensed much further, to what they concluded or produced — as long as nothing in the six items above is dropped.
Create a complete Markdown handoff file from this conversation.”*

New Chat Start-Off Prompt

Paste into New Chat Thread:

“Use the attached handoff file as the working record for this matter. First list the objective, the verified facts, open questions, constraints, and the next recommended action. Then wait for my instruction. Do not add facts not supported by the handoff file or any attached source.”

Document Parsing Prompt (before analyzing any Large Doc)

You are reviewing the attached document for [business purpose]. Do not analyze or summarize the entire document. Create a document map in a table with section or page range; topic; key facts; decision relevance to [purpose]; exact quotes worth preserving; whether the section is germane, meaning directly relevant, to the purpose.

CHART and TABLE CREATION PROMPTING

If you are trying to create an image from a data set (a visual chart or table), you will find that may often create an image with errors in the numbers. One Pro Tip is to prompt an AI to Code the HTML Image for you. Then upload that image into a design tool like Canva, Adobe etc

Upload an image sample you like and prompt this:

“Code an HTML infographic like the one I attached, but about (topic) using these numbers (include a CSV or paste numbers you want)”

The text is spelled correctly. Image generators still can't promise that.

Part Five: If You Slip

It happens. What you do in the next ten minutes decides whether it stays a mistake.

- Stop using the thread. Do not delete it. Deleting destroys the record your privacy office needs, and deleted chats have been ordered produced as evidence anyway.
- Note what was pasted, into which tool, at what time, and on whose account.
- Notify your manager and the person responsible for IT security or data protection immediately. Do not wait until the end of the day.
- Let the privacy office make the legal calls on notification, risk assessment, and mitigation. Your job is to surface it fast and completely.

Reporting a mistake promptly is the action that keeps it a mistake. Hiding it is the action that makes it a headline.

AI Privacy Checklist

01 Choose Right Plan Use Enterprise or Sandboxed tool plan for work. Check seat minimums and data policies of vendors-confirm BAA and independent compliance verification.	02 Turn OFF Training & Verify Retention Disable model training and check retention rules of chosen plan since NDAs and customer agreements may be violated.
03 Skip Chat Feedback Avoid thumb up/down on sensitive chats. Feedback may share conversation violating compliance or NDAs or customer agreements.	04 Keep Chats Private Do Not Create Public Share Links for sensitive chats.
05 Use Temp Chat or Incognito For sensitive topics. Safety retention may still apply.	06 Turn Off Memory Disable memory for more privacy. Review/delete old memories.
07 Delete Old Links Prior shared chats can be deleted to avoid future compliance violations.	08 Check Codex Settings Review separate Codex training controls, including full-environment data.
09 Verify Browser Extensions Check who can read chats and where data gets shared.	10 Protect Sensitive Data Keep PHI, legal, financial, client and unreleased IP out of all personal accounts and verify retention rules.

Quick Reference Card

Always

- Sanitize before you type, not after.
- Use only company-approved tools for work. Do not include confidential data unless the tool and use are explicitly authorized. Substitute roles and ranges instead of deleting facts.
- Add Role, Examples, and eXpectations to every real task.
- End prompts with "Ask me questions as needed before answering." Keep the Verification Prompt in global instructions.
- Fresh chat every 15 prompts.

Never

- Names, full dates, Social Security numbers, customer or employee IDs, addresses, phone numbers, or email addresses.
- Confidential audio, face photos, or credentials. In any form.
- The original file when a paragraph will do.
- Give clients copy-paste AI without verification and Human review and editing..
- A document approved or sent without checking it against the original source and what actually happened.

Want to learn more

Ask about our complimentary AI Workshop for leadership, or our complimentary Live Cyber Awareness Training for all employees. Please share this with others.

A sanitized prompt is not a weaker prompt. It is the same question, minus the liability.

Created by David Dean Mauro | NetGainIT.com | CyberCrimeJunkies.com | Moving Target Book Series